

# Accessing the contents of a stanfit object

*Stan Development Team*

---

**2019-07-09**

---

- [Posterior draws](#)
- [Posterior summary statistics and convergence diagnostics](#)
- [Sampler diagnostics](#)
- [Model code](#)
- [Initial values](#)
- [\(P\)RNG seed](#)
- [Warmup and sampling times](#)

This vignette demonstrates how to access most of data stored in a stanfit object. A stanfit object (an object of class "stanfit") contains the output derived from fitting a Stan model using Markov chain Monte Carlo or one of Stan's variational approximations (meanfield or full-rank). Throughout the document we'll use the stanfit object obtained from fitting the Eight Schools example model:

```
library(rstan)
fit <- stan_demo("eight_schools", refresh = 0)
```

```
class(fit)
```

```
[1] "stanfit"
attr(,"package")
[1] "rstan"
```

## Posterior draws

---

There are several functions that can be used to access the draws from the posterior distribution stored in a stanfit object. These are `extract`, `as.matrix`, `as.data.frame`, and `as.array`, each of which returns the draws in a different format.

### `extract()`

---

The `extract` function (with its default arguments) returns a list with named components corresponding to the model parameters.

```
list_of_draws <- extract(fit)
```

```
print(names(list_of_draws))
```

```
[1] "mu"      "tau"      "eta"      "theta" "lp__"
```

In this model the parameters `mu` and `tau` are scalars and `theta` is a vector with eight elements. This means that the draws for `mu` and `tau` will be vectors (with length equal to the number of post-warmup iterations times the number of chains) and the draws for `theta` will be a matrix, with each column corresponding to one of the eight components:

```
head(list_of_draws$mu)
```

```
[1] 12.252686 14.907965 -1.373670  5.154424  6.001779 16.319632
```

```
head(list_of_draws$tau)
```

```
[1] 13.510971 10.187284  7.714862 10.243666 16.823645 17.440304
```

```
head(list_of_draws$theta)
```

```
iterations      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
[1,] 15.60367  7.588454  1.576459 19.651057  7.664926  1.0059249
[2,] 15.34390  5.364125 20.853174 14.643021  7.335447  5.0269569
[3,] 12.34758  8.180855  7.766418  5.044026  8.440089  7.1604064
[4,] 19.65023 -4.479000 -6.509202 -4.066218 -5.015240  3.7905639
[5,] 38.23529 -5.416479 -2.989955  6.736584 19.132428  0.8758841
[6,] 12.03491 11.660757 19.737807 16.787683  7.368761  5.8037854
```

```
iterations      [,7]      [,8]
[1,]  5.284314 15.774425
[2,] 29.245463 21.365687
[3,]  8.908072 -11.216071
[4,]  2.958357  8.495162
[5,]  8.858071 23.496445
[6,] 19.364599 25.048120
```

**as.matrix(), as.data.frame(), as.array()**

The `as.matrix`, `as.data.frame`, and `as.array` functions can also be used to retrieve the posterior draws from a stanfit object:

```
matrix_of_draws <- as.matrix(fit)
```

```
print(colnames(matrix_of_draws))
```

```
[1] "mu"      "tau"      "eta[1]"   "eta[2]"   "eta[3]"   "eta[4]"
[7] "eta[5]"   "eta[6]"   "eta[7]"   "eta[8]"   "theta[1]" "theta[2]"
[13] "theta[3]" "theta[4]" "theta[5]" "theta[6]" "theta[7]" "theta[8]"
[19] "lp__"
```

```
df_of_draws <- as.data.frame(fit)
print(colnames(df_of_draws))
```

```
[1] "mu"      "tau"      "eta[1]"   "eta[2]"   "eta[3]"   "eta[4]"
[7] "eta[5]"   "eta[6]"   "eta[7]"   "eta[8]"   "theta[1]" "theta[2]"
[13] "theta[3]" "theta[4]" "theta[5]" "theta[6]" "theta[7]" "theta[8]"
[19] "lp__"
```

```
array_of_draws <- as.array(fit)
print(dimnames(array_of_draws))
```

```
$iterations
NULL
```

```
$chains
[1] "chain:1" "chain:2" "chain:3" "chain:4"
```

```
$parameters
[1] "mu"      "tau"      "eta[1]"   "eta[2]"   "eta[3]"   "eta[4]"
[7] "eta[5]"   "eta[6]"   "eta[7]"   "eta[8]"   "theta[1]" "theta[2]"
[13] "theta[3]" "theta[4]" "theta[5]" "theta[6]" "theta[7]" "theta[8]"
[19] "lp__"
```

The `as.matrix` and `as.data.frame` methods essentially return the same thing except in matrix and data frame form, respectively. The `as.array` method returns the draws from each chain separately and so has an additional dimension:

```
print(dim(matrix_of_draws))
print(dim(df_of_draws))
print(dim(array_of_draws))
```

```
[1] 4000  19
[1] 4000  19
[1] 1000   4  19
```

By default all of the functions for retrieving the posterior draws return the draws for *all* parameters (and generated quantities). The optional argument `pars` (a character vector) can be used if only a subset of

the parameters is desired, for example:

```
mu_and_theta1 <- as.matrix(fit, pars = c("mu", "theta[1]"))
head(mu_and_theta1)
```

```
      parameters
iterations      mu  theta[1]
[1,] 11.594856 14.4775138
[2,] 14.105900 28.7191144
[3,]  2.790656  3.5523486
[4,]  7.954872  0.1407352
[5,]  6.731532  9.2016480
[6,]  1.242287 31.3440051
```

## Posterior summary statistics and convergence diagnostics

Summary statistics are obtained using the `summary` function. The object returned is a list with two components:

```
fit_summary <- summary(fit)
print(names(fit_summary))
```

```
[1] "summary"  "c_summary"
```

In `fit_summary$summary` all chains are merged whereas `fit_summary$c_summary` contains summaries for each chain individually. Typically we want the summary for all chains merged, which is what we'll focus on here.

The summary is a matrix with rows corresponding to parameters and columns to the various summary quantities. These include the posterior mean, the posterior standard deviation, and various quantiles computed from the draws. The `probs` argument can be used to specify which quantiles to compute and `pars` can be used to specify a subset of parameters to include in the summary.

For models fit using MCMC, also included in the summary are the Monte Carlo standard error (`se_mean`), the effective sample size (`n_eff`), and the R-hat statistic (`Rhat`).

```
print(fit_summary$summary)
```

```
      mean      se_mean      sd      2.5%      25%
mu      8.40239654 0.23939862 5.6477978 -1.5856605  4.9203079
tau      6.81657128 0.24661168 6.0553821  0.1817417  2.5737606
eta[1]    0.38782960 0.01503165 0.9163859 -1.4184593 -0.2161574
eta[2]   -0.01614640 0.01421635 0.8720442 -1.6918910 -0.6043018
eta[3]   -0.23521562 0.01768225 0.9285214 -1.9991309 -0.8742552
eta[4]   -0.04992150 0.01389894 0.8771066 -1.7692245 -0.6322952
```

```

eta[5]      -0.37046390 0.01595938 0.8976414 -2.0454028 -0.9840520
eta[6]      -0.23426162 0.01543708 0.8877458 -1.9427380 -0.8319003
eta[7]       0.34358924 0.01626435 0.8855897 -1.4118509 -0.2502239
eta[8]       0.03774424 0.01473862 0.9480062 -1.8070263 -0.5896104
theta[1]    11.67308393 0.18109835 8.3804917 -1.6045395 6.1696642
theta[2]     8.07366335 0.09534171 6.2302621 -3.8610079 4.0153215
theta[3]     5.99503306 0.14031161 7.8550471 -12.0653111 1.8476810
theta[4]     7.82824007 0.11181224 6.5417373 -5.0297220 3.7450550
theta[5]     5.09580003 0.10321987 6.5831540 -9.4541310 1.2067127
theta[6]     6.28507381 0.11712779 6.8298216 -8.4880139 2.3255196
theta[7]    10.86160511 0.15631508 6.8499727 -0.9571675 6.2709256
theta[8]     8.58388560 0.19459084 8.0906663 -6.6119231 3.6710090
lp__        -39.55676603 0.07746429 2.7015009 -45.5703000 -41.1940633

```

	50%	75%	97.5%	n_eff	Rhat
mu	8.20833569	11.5279047	20.510302	556.5638	1.0051426
tau	5.35894999	9.3390741	22.420817	602.9146	1.0089599
eta[1]	0.39503008	1.0036727	2.199559	3716.5797	1.0001790
eta[2]	-0.03301239	0.5607387	1.736386	3762.7118	0.9998571
eta[3]	-0.25034213	0.3901198	1.616674	2757.4571	1.0015253
eta[4]	-0.05807728	0.5077293	1.711333	3982.3682	1.0016549
eta[5]	-0.39406939	0.2151956	1.474333	3163.5424	0.9994128
eta[6]	-0.23986139	0.3256300	1.576667	3307.0963	1.0002083
eta[7]	0.33572257	0.9322493	2.075483	2964.7765	0.9998975
eta[8]	0.03800369	0.6691663	1.912977	4137.2222	0.9995512
theta[1]	10.52117034	15.6999619	32.100103	2141.4602	1.0024168
theta[2]	8.05066651	11.9398324	21.131261	4270.1858	0.9993536
theta[3]	6.59203659	10.8666143	19.738102	3134.0820	0.9996507
theta[4]	7.84118702	11.7344730	21.401794	3423.0049	0.9999652
theta[5]	5.57323870	9.4067905	16.804645	4067.6299	0.9997760
theta[6]	6.43362090	10.5563822	19.225530	3400.1558	1.0011304
theta[7]	10.31091633	14.6990088	26.365838	1920.3302	1.0029145
theta[8]	8.24268197	12.8361120	25.946117	1728.7166	1.0018890
lp__	-39.24268861	-37.6783486	-34.997607	1216.2059	1.0008591

If, for example, we wanted the only quantiles included to be 10% and 90%, and for only the parameters included to be `mu` and `tau`, we would specify that like this:

```

mu_tau_summary <- summary(fit, pars = c("mu", "tau"), probs = c(0.1, 0.9))$summary
print(mu_tau_summary)

```

```

      mean  se_mean      sd      10%      90%  n_eff  Rhat
mu  8.402397 0.2393986 5.647798 1.8414735 14.75461 556.5638 1.005143
tau 6.816571 0.2466117 6.055382 0.9959847 14.28546 602.9146 1.008960

```

Since `mu_tau_summary` is a matrix we can pull out columns using their names:

```

mu_tau_80pct <- mu_tau_summary[, c("10%", "90%")]
print(mu_tau_80pct)

```

	10%	90%
mu	1.8414735	14.75461
tau	0.9959847	14.28546

## Sampler diagnostics

For models fit using MCMC the stanfit object will also contain the values of parameters used for the sampler. The `get_sampler_params` function can be used to access this information.

The object returned by `get_sampler_params` is a list with one component (a matrix) per chain. Each of the matrices has number of columns corresponding to the number of sampler parameters and the column names provide the parameter names. The optional argument `inc_warmup` (defaulting to `TRUE`) indicates whether to include the warmup period.

```
sampler_params <- get_sampler_params(fit, inc_warmup = FALSE)
sampler_params_chain1 <- sampler_params[[1]]
colnames(sampler_params_chain1)
```

```
[1] "accept_stat__" "stepsize__"    "treedepth__"   "n_leapfrog__"
[5] "divergent__"   "energy__"
```

To do things like calculate the average value of `accept_stat__` for each chain (or the maximum value of `treedepth__` for each chain if using the NUTS algorithm, etc.) the `sapply` function is useful as it will apply the same function to each component of `sampler_params`:

```
mean_accept_stat_by_chain <- sapply(sampler_params, function(x) mean(x[,
"accept_stat__"]))
print(mean_accept_stat_by_chain)
```

```
[1] 0.7945660 0.8666220 0.9066320 0.8979723
```

```
max_treedepth_by_chain <- sapply(sampler_params, function(x) max(x[, "treedepth__"]))
print(max_treedepth_by_chain)
```

```
[1] 5 5 5 5
```

## Model code

The Stan program itself is also stored in the stanfit object and can be accessed using `get_stancode`:

```
code <- get_stancode(fit)
```

The object code is a single string and is not very intelligible when printed:

```
print(code)
```

```
[1] "data {\n  int<lower=0> J;          // number of schools \n  real y[J];\n  // estimated treatment effects\n  real<lower=0> sigma[J]; // s.e. of effect estimates\n}\n\nparameters {\n  real mu; \n  real<lower=0> tau;\n  vector[J] eta;\n}\n\ntransformed\nparameters {\n  vector[J] theta;\n  theta = mu + tau * eta;\n}\n\nmodel {\n  target +=\n  normal_lpdf(eta | 0, 1);\n  target += normal_lpdf(y | theta, sigma);\n}"
attr(,"model_name2")
[1] "schools"
```

A readable version can be printed using cat:

```
cat(code)
```

```
data {
  int<lower=0> J;          // number of schools
  real y[J];             // estimated treatment effects
  real<lower=0> sigma[J]; // s.e. of effect estimates
}
parameters {
  real mu;
  real<lower=0> tau;
  vector[J] eta;
}
transformed parameters {
  vector[J] theta;
  theta = mu + tau * eta;
}
model {
  target += normal_lpdf(eta | 0, 1);
  target += normal_lpdf(y | theta, sigma);
}
```

## Initial values

The `get_inits` function returns initial values as a list with one component per chain. Each component is itself a (named) list containing the initial values for each parameter for the corresponding chain:

```
inits <- get_inits(fit)
inits_chain1 <- inits[[1]]
```

```
print(inits_chain1)
```

```
$mu  
[1] -0.201507
```

```
$tau  
[1] 0.1682838
```

```
$eta  
[1] -1.6057067 -1.8569811 0.3266007 1.0876453 -0.7585606 -0.4009873  
[7] -1.7051884 -0.8772185
```

```
$theta  
[1] -0.47172139 -0.51400680 -0.14654537 -0.01847388 -0.32916042 -0.26898663  
[7] -0.48846255 -0.34912864
```

## (P)RNG seed

---

The `get_seed` function returns the (P)RNG seed as an integer:

```
print(get_seed(fit))
```

```
[1] 886183407
```

## Warmup and sampling times

---

The `get_elapsed_time` function returns a matrix with the warmup and sampling times for each chain:

```
print(get_elapsed_time(fit))
```

```
      warmup  sample  
chain:1 0.059150 0.034474  
chain:2 0.036554 0.035265  
chain:3 0.039221 0.040196  
chain:4 0.036331 0.039571
```