

Bayesian modelling using R + STAN = RSTSAN

https://github.com/clobos/Seminario_STAN_UNESP

Cristian Villegas, ESALQ/USP

UNESP Botucatu (29/08/2019)

Stan



Brief introduction to Stan

Stan is a state-of-the-art platform for statistical modeling and high-performance statistical computation. Thousands of users rely on Stan for statistical modeling, data analysis, and prediction in the social, biological, and physical sciences, engineering, and business.

Brief introduction to Stan

Users specify log density functions in Stan's probabilistic programming language and get:

- ▶ full Bayesian statistical inference with MCMC sampling (NUTS, HMC)
- ▶ approximate Bayesian inference with variational inference (ADVI)
- ▶ penalized maximum likelihood estimation with optimization (L-BFGS)

Brief introduction to Stan

Stan's math library provides differentiable probability functions & linear algebra (C++ autodiff). Additional R packages provide expression-based linear modeling, posterior visualization, and leave-one-out cross-validation.

Motivation Example




Figure 1: Scatterplot of length-at-capture (cm) versus age-at-capture (years)

The Von Bertalanffy growth model (Classical approach)

The Von Bertalanffy growth model is given by

$$Y_i = \beta_1 (1 - \exp\{-\beta_2 (x_i - \beta_3)\}) + \varepsilon_i \quad \text{with} \quad \varepsilon_i \sim N(0, \sigma^2),$$

where Y_i is the length-at-capture and x_i is the age-at-capture for the i th fish, respectively. The parameter interpretation of the Von Bertalanffy growth model is given by

- ▶ β_1 is the asymptotic (average) length
- ▶ β_2 is the growth rate coefficient (units are years^{-1}), and
- ▶ β_3 represent the age when (average) length was zero.

The Von Bertalanffy growth model (Classical approach)

In R code the Von Bertalanffy growth model is given by

```
LVB<-function(x,beta1, beta2, beta3){  
  beta1*(1-exp(-beta2*(x-beta3)))  
}
```


The Von Bertalanffy growth model (Classical approach)

```
fit_LVB <- nls(Length ~ LVB(Age, beta1, beta2, beta3),  
start = list(beta1 = max(dados.amostra$Length),  
              beta2 = 0.5,  
              beta3 = 0),  
data = dados.amostra)  
  
c(coef(fit_LVB), sigma=sigma(fit_LVB))
```

```
##          beta1          beta2          beta3          sigma  
## 73.8227693    0.1080882   -3.1551711    3.2475942
```

First Case: Bayesian Inference with Stan code

```
Von_Bertanlanffy_mcmc <- '  
data {  
  int<lower = 0> N ;  
  vector[N] x ;  
  vector[N] y ;  
}  
parameters {  
  real<lower = .0> beta1 ;  
  real<lower = .0> beta2 ;  
  real beta3 ;  
  real<lower = .0> sigma ;  
}  
model {  
  y ~ normal(beta1*(1-exp(-beta2*(x-beta3))), sigma) ;  
}  
'
```

Bayesian Inference: Stan code

```
fit_Von_Bertanlanffy <- stan(  
  model_code = Von_Bertanlanffy_mcmc,  
  data = list(N = nrow(dados.amostra),  
              x = dados.amostra$Age,  
              y = dados.amostra$Length),  
  chain = 3,  
  iter = 100,  
  warmup = 10,  
  thin = 1,  
  refresh=0)
```

MCMC diagnostics using the bayesplot package

```
parameters<- c(paste('beta',1:3, sep=""), 'sigma')  
  
CI_theta <- summary(ajuste_Von_Bertanlanffy,  
                    pars = parameters,  
                    probs = c(0.025, 0.975))$summary  
print(round(CI_theta,2))
```

##		mean	se_mean	sd	2.5%	97.5%	n_eff	Rhat
##	beta1	25.45	13.46	21.95	0.37	50.18	2.66	2.30
##	beta2	2.46	1.88	2.31	0.54	5.70	1.52	17.00
##	beta3	0.84	0.30	0.45	0.35	1.91	2.28	1.97
##	sigma	10.32	5.17	12.87	0.90	42.48	6.20	1.78

MCMC diagnostics using the bayesplot package

```
traceplot(ajuste_Von_Bertanlanffy, pars = parameters,  
          inc_warmup = TRUE)
```

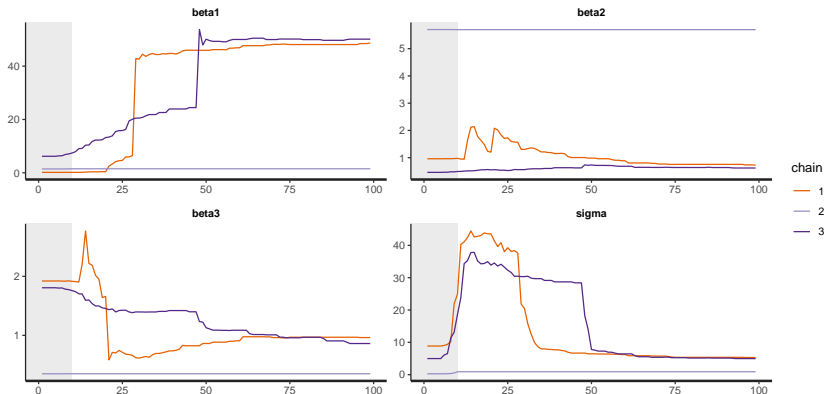


Figure 2: Traceplots for the Von Bertalanffy growth model

MCMC diagnostics using the bayesplot package

```
mcmc_combo(mcmc_chain, pars = parameters, n_warmup=0)
```

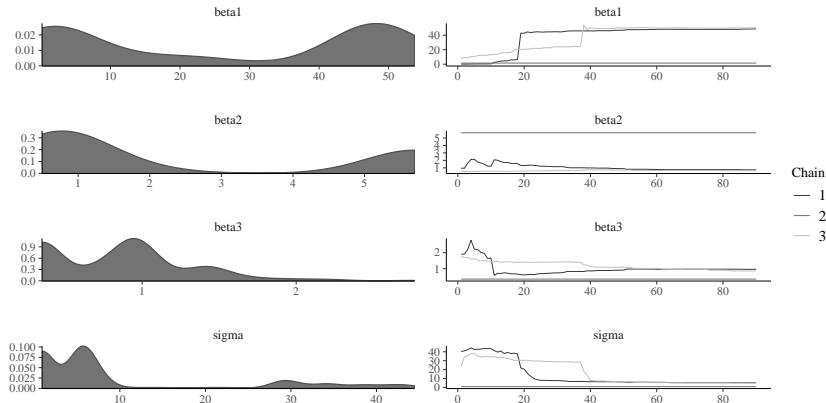


Figure 3: Posterior distributions and traceplots for the Von Bertalanffy growth model

MCMC diagnostics using the bayesplot package

```
mcmc_pairs(mcmc_chain, pars = parameters)
```

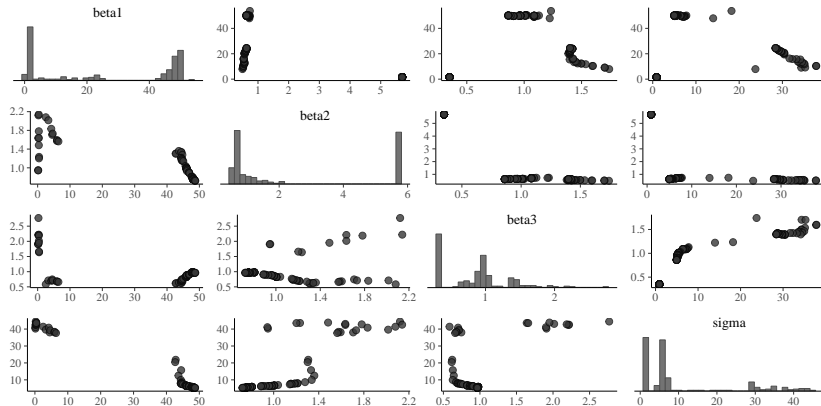


Figure 4: Scatterplots of MCMC draws for the Von Bertalanffy growth model

Comparing fitted curves based on Classical and Bayesian Inference

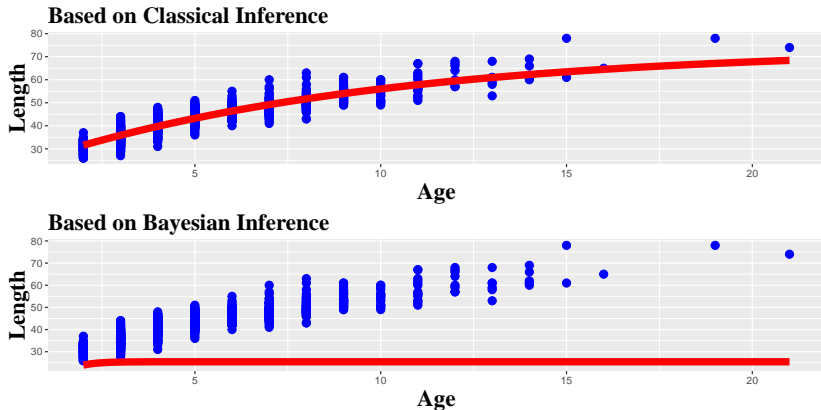


Figure 5: Fitted curves based on classical and bayesian Inference for the Von Bertalanffy growth model

Comparing the parameter estimation based on Classical and Bayesian Inference

```
c(coef(fit_LVB),sigma=sigma(fit_LVB))
```

```
##          beta1          beta2          beta3          sigma
## 73.8227693   0.1080882 -3.1551711   3.2475942
```

```
print(CI_theta[,1])
```

```
##          beta1          beta2          beta3          sigma
## 25.4525584   2.4567863   0.8429762 10.3197269
```

Second Case: Bayesian Inference with Stan code

```
Von_Bertanlanffy_mcmc <- '  
data {  
  int<lower = 0> N ;  
  vector[N] x ;  
  vector[N] y ;  
}  
parameters {  
  real<lower = .0> beta1 ;  
  real<lower = .0> beta2 ;  
  real beta3 ;  
  real<lower = .0> sigma ;  
}  
model {  
  y ~ normal(beta1*(1-exp(-beta2*(x-beta3))), sigma) ;  
}  
'
```

Bayesian Inference: Stan code

```
fit_Von_Bertanlanffy <- stan(  
  model_code = Von_Bertanlanffy_mcmc,  
  data = list(N = nrow(dados.amostra),  
              x = dados.amostra$Age,  
              y = dados.amostra$Length),  
  chain = 3,  
  iter = 11000,  
  warmup = 1000,  
  thin = 10,  
  refresh=0)
```

MCMC diagnostics using the bayesplot package

```
parameters<- c(paste('beta',1:3, sep=""), 'sigma')  
  
CI_theta <- summary(ajuste_Von_Bertanlanffy,  
                    pars = parameters,  
                    probs = c(0.025, 0.975))$summary  
print(round(CI_theta,2))
```

##		mean	se_mean	sd	2.5%	97.5%	n_eff	Rhat
##	beta1	74.09	0.05	2.54	69.63	79.49	2847.90	1
##	beta2	0.11	0.00	0.01	0.09	0.13	2885.81	1
##	beta3	-3.18	0.01	0.28	-3.76	-2.66	2932.61	1
##	sigma	3.25	0.00	0.07	3.11	3.40	3051.84	1

MCMC diagnostics using the bayesplot package

```
traceplot(ajuste_Von_Bertanlanffy, pars = parameters,  
          inc_warmup = TRUE)
```

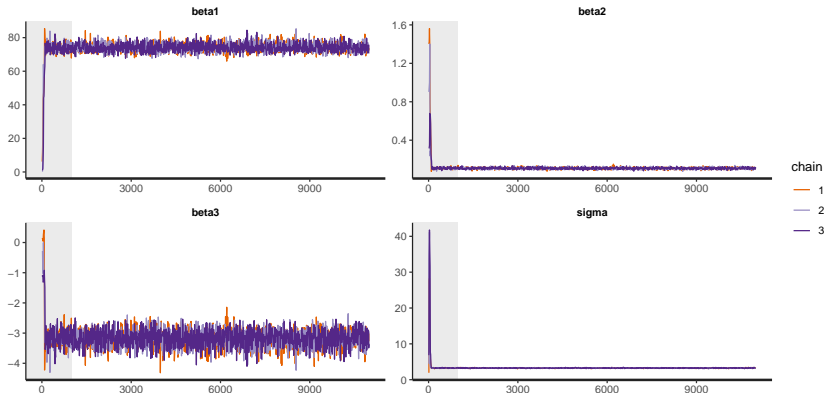


Figure 6: Traceplots for the Von Bertalanffy growth model

MCMC diagnostics using the bayesplot package

```
mcmc_combo(mcmc_chain, pars = parameters, n_warmup=0)
```

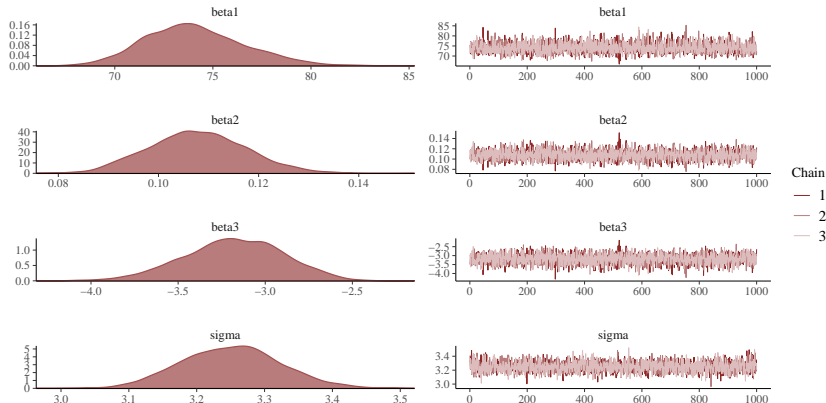


Figure 7: Posterior distributions and traceplots for the Von Bertalanffy growth model

MCMC diagnostics using the bayesplot package

```
mcmc_pairs(mcmc_chain, pars = parameters)
```




Figure 8: Scatterplots of MCMC draws for the Von Bertalanffy growth model

Comparing fitted curves based on Classical and Bayesian Inference

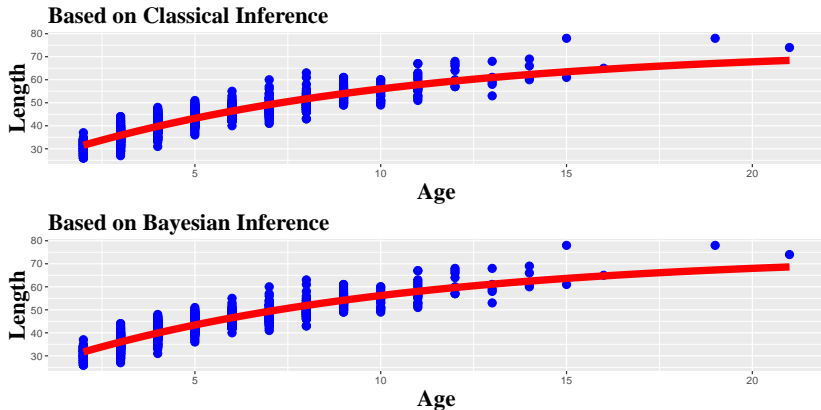


Figure 9: Fitted curves based on classical and bayesian Inference for the Von Bertalanffy growth model

Comparing the parameter estimation based on Classical and Bayesian Inference

```
c(coef(fit_LVB),sigma=sigma(fit_LVB))
```

```
##          beta1          beta2          beta3          sigma
## 73.8227693  0.1080882 -3.1551711  3.2475942
```

```
print(CI_theta[,1])
```

```
##          beta1          beta2          beta3          sigma
## 74.0883836  0.1077937 -3.1803626  3.2509435
```

Normal distribution without covariates in Stan

$Y \sim N(\mu, \sigma^2)$. Therefore, $\theta = (\mu, \sigma^2)^\top$. Here, we specify the prior distribution for each parameter.

Bayesian Inference: Stan code (specific prior)

```
normal_dist_example<- '  
data {  
  int<lower=0> N;  
  vector[N] y;  
}  
parameters {  
  real mu;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(mu, sigma);  
  mu~ normal(0,1e6);  
  sigma ~ student_t(3,0,1);  
}  
,
```

Bayesian Inference: Stan code (specific prior)

```
normal_dist_fit <- stan(model_code = normal_dist_example,  
  data = list(N = dim(stackloss)[1],  
    y = stackloss$stack.loss),  
  chain = 3,  
  iter = 11000,  
  warmup = 1000,  
  thin = 10,  
  refresh = 0)
```

MCMC diagnostics using the bayesplot package

```
parameters<- c("mu", "sigma")

CI_theta <- summary(normal_dist_fit,
                    pars = parameters,
                    probs = c(0.025, 0.975))$summary
print(round(CI_theta,2))
```

##	mean	se_mean	sd	2.5%	97.5%	n_eff	Rhat
## mu	17.50	0.04	2.15	13.18	21.80	2970.46	1
## sigma	9.82	0.03	1.53	7.40	13.27	2973.68	1

MCMC diagnostics using the bayesplot package

```
traceplot(normal_dist_fit, pars = parameters,  
          inc_warmup = TRUE)
```

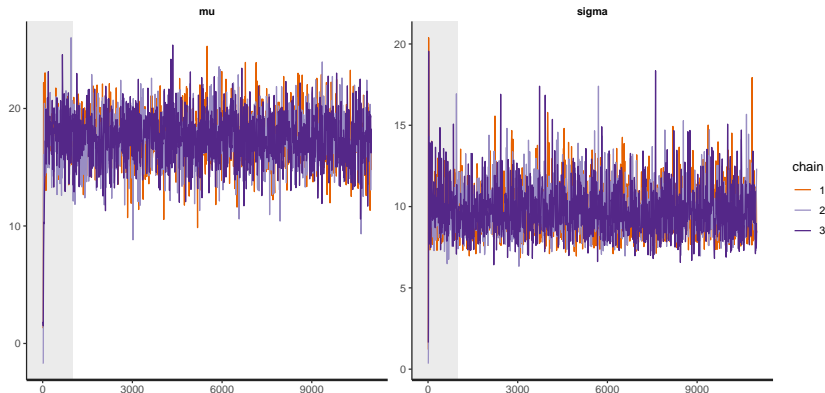


Figure 10: Traceplots for the Normal distribution

MCMC diagnostics using the bayesplot package

```
mcmc_combo(mcmc_chain, pars = parameters, n_warmup=0)
```

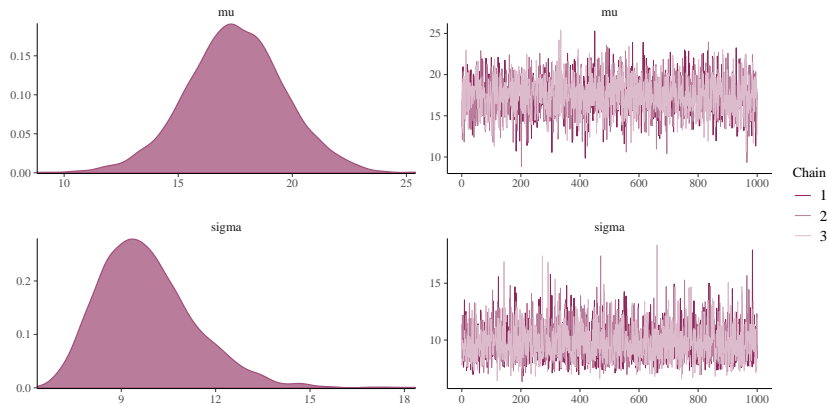


Figure 11: Posterior distributions and traceplots for the Normal distribution

MCMC diagnostics using the bayesplot package

```
mcmc_pairs(mcmc_chain, pars = parameters)
```




Figure 12: Scatterplots of MCMC draws for the Normal distribution

Motivation Example




Figure 13: Scatterplot of proportions versus $\log(\text{Dose})$

Logistic Regression

These are the number of adult flour beetles which died following a 5-hour exposure to gaseous carbon disulphide. WBinomial response with logit link function. Here, we do not specify the prior distribution for each parameter.

Bayesian Inference: Stan code

```
logistic_example<- 'data {  
  int<lower=0> N;vector[N] x;  
  int<lower=0> y[N]; int<lower=0> n[N];  
}  
parameters {real beta1;real beta2;  
}  
transformed parameters {  
  real exp_eta[N]; real<lower=0, upper=1> prob[N];  
  for (i in 1:N) {exp_eta[i] = exp(beta1 + beta2*x[i]);  
  prob[i]= exp_eta[i]/(exp_eta[i] + 1);  
}  
}  
model {  
  y ~ binomial_logit(n, beta1 + beta2 * x);  
}  
,
```

Bayesian Inference: Stan code

```
logistic_fit <- stan(model_code = logistic_example,  
data = list(N = dim(beetleDat)[1],  
             n = beetleDat$n,  
             x = beetleDat$lDose,  
             y = beetleDat$x),  
             chain = 3,  
             iter = 11000,  
             warmup = 1000,  
             thin = 10,  
             refresh=0)
```

MCMC diagnostics using the bayesplot package

```
parameters<- c(paste('beta',1:2, sep=""))  
  
CI_theta <- summary(logistic_fit,  
                    pars = parameters,  
                    probs = c(0.025, 0.975))$summary  
print(round(CI_theta,2))
```

##		mean	se_mean	sd	2.5%	97.5%	n_eff	Rhat
##	beta1	-61.37	0.11	5.25	-72.01	-51.43	2420.70	1
##	beta2	34.64	0.06	2.95	29.07	40.58	2422.88	1

MCMC diagnostics using the bayesplot package

```
traceplot(logistic_fit, pars = parameters,  
          inc_warmup = TRUE)
```

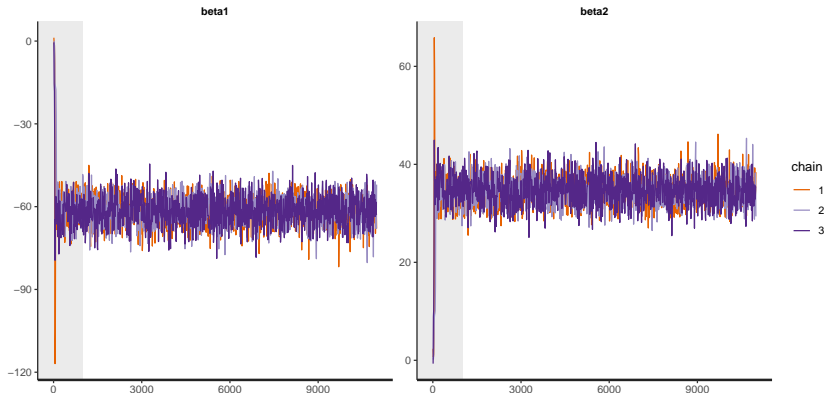


Figure 14: Traceplots for the Logistic regression model

MCMC diagnostics using the bayesplot package

```
mcmc_combo(mcmc_chain, pars = parameters, n_warmup=0)
```

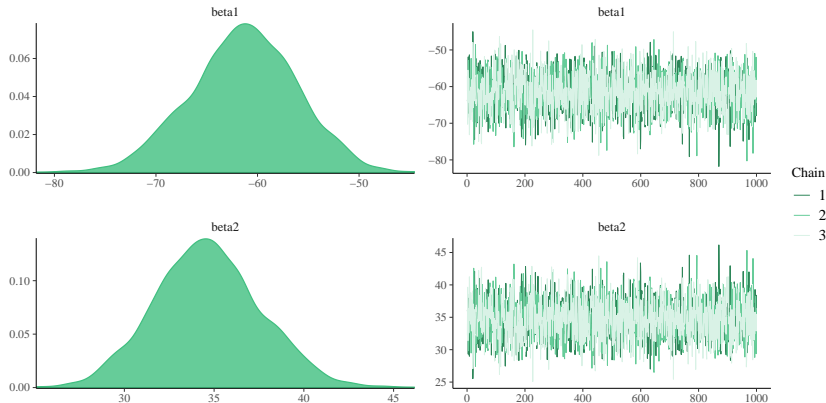


Figure 15: Posterior distributions and traceplots for the Logistic regression model

MCMC diagnostics using the bayesplot package

```
mcmc_pairs(mcmc_chain, pars = parameters)
```




Figure 16: Scatterplots of MCMC draws for the Logistic Regression model

<https://www.youtube.com/watch?v=uSjsJg8fcwY>

Fitted curve based on Bayesian Inference

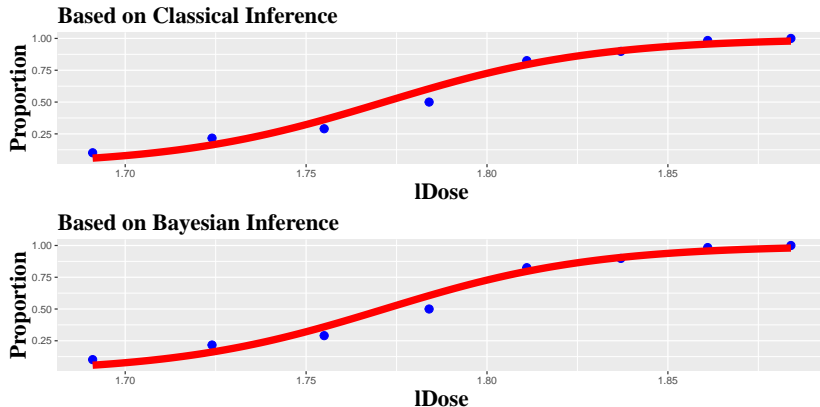


Figure 17: Fitted curves based on classical and bayesian Inference for the Logistic Regression model

Comparing the parameter estimation based on Classical and Bayesian Inference

```
c(beta1=coef(fitLogistic)[1],beta2=coef(fitLogistic)[2])
```

```
## beta1.(Intercept)      beta2.lDose  
##           -60.74013           34.28593
```

```
print(CI_theta[,1])
```

```
##      beta1      beta2  
## -61.37130  34.64121
```

Motivation Example




Figure 18: Scatterplot of stack.loss versus Air.Flow

Normal Linear Model

$$Y_i = \beta_1 + \beta_2 x_i + \varepsilon_i \quad \text{with} \quad \varepsilon_i \sim N(0, \sigma^2),$$

where Y_i is stack loss and x_i is flow of cooling air, respectively. Therefore $\theta = (\beta_1, \beta_2, \sigma^2)^\top$. Here, we do not specify the prior distribution for each parameter.

Bayesian Inference: Stan code

```
lm_example<- '  
data {  
  int<lower=0> N;  
  vector[N] x;  
  vector[N] y;  
}  
parameters {  
  real beta1;  
  real beta2;  
  real<lower=0> sigma;  
}  
model {  
  y ~ normal(beta1+beta2*x, sigma);  
}  
'
```

Bayesian Inference: Stan code

```
lm_fit <- stan(model_code = lm_example,  
data = list(N = dim(stackloss)[1],  
            x = stackloss$Air.Flow,  
y = stackloss$stack.loss),  
            chain = 3,  
            iter = 11000,  
            warmup = 1000,  
            thin = 10,  
            refresh=0)
```

MCMC diagnostics using the bayesplot package

```
parameters<- c(paste('beta',1:2, sep=""), 'sigma')

CI_theta <- summary(lm_fit,
                    pars = parameters,
                    probs = c(0.025, 0.975))$summary
print(round(CI_theta,2))
```

##		mean	se_mean	sd	2.5%	97.5%	n_eff	Rhat
##	beta1	-43.85	0.12	6.57	-56.83	-30.91	2816.31	1
##	beta2	1.02	0.00	0.11	0.80	1.23	2831.99	1
##	sigma	4.39	0.01	0.77	3.20	6.14	2976.60	1

MCMC diagnostics using the bayesplot package

```
traceplot(lm_fit, pars = parameters, inc_warmup = TRUE)
```




Figure 19: Traceplots for the Linear Normal

MCMC diagnostics using the bayesplot package

```
mcmc_combo(mcmc_chain, pars = parameters, n_warmup=0)
```

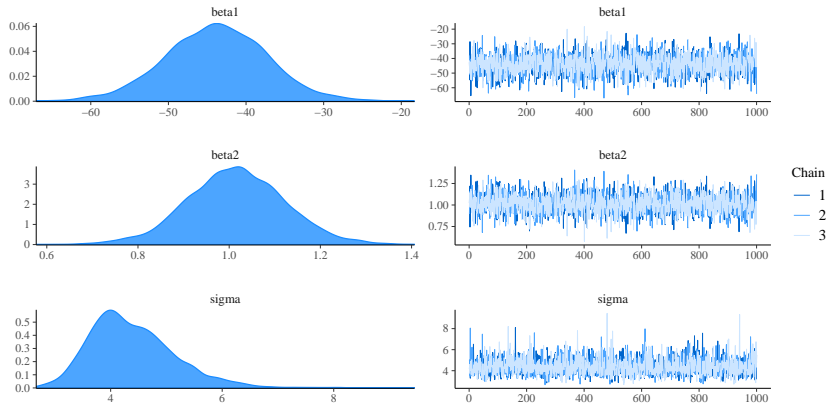


Figure 20: Posterior distributions and traceplots for the Linear Normal model

MCMC diagnostics using the bayesplot package

```
mcmc_pairs(mcmc_chain, pars = parameters)
```

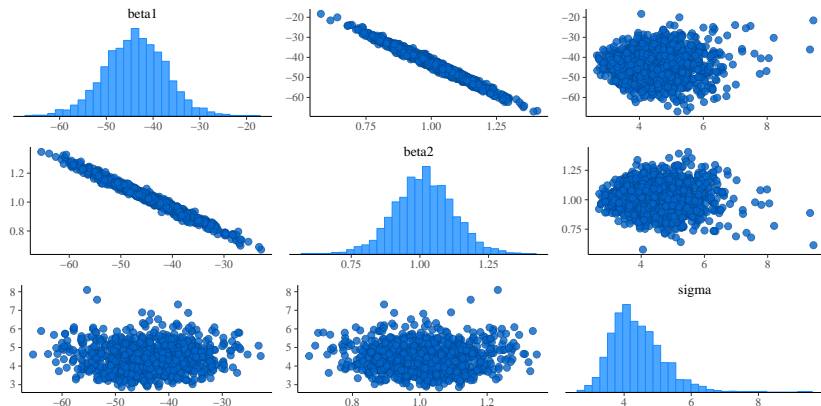


Figure 21: Scatterplots of MCMC draws for the Linear Normal model

Comparing the parameter estimation based on Classical and Bayesian Inference

```
c(beta1=coef(fit)[1],beta2=coef(fit)[2],  
  sigma=sigma(fit))
```

##	beta1.(Intercept)	beta2.Air.Flow	sigma
##	-44.132025	1.020309	4.098242

```
print(CI_theta[,1])
```

##	beta1	beta2	sigma
##	-43.847882	1.015100	4.393619

Comparing fitted curves based on Classical and Bayesian Inference

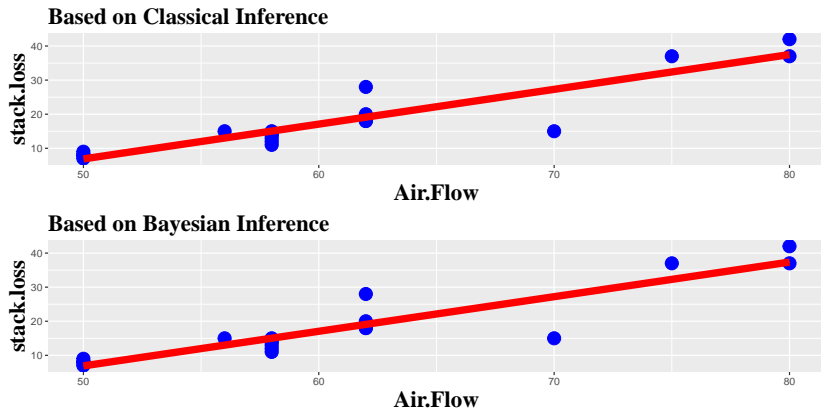


Figure 22: Fitted curves based on classical and bayesian Inference for the Normal Linear model

More R packages based on Stan

- ▶ Bayesian Applied Regression Modeling via Stan: `rstanarm` R package.
- ▶ Interactive Visual and Numerical Diagnostics and Posterior Analysis for Bayesian Models `shinystan` R package.

Do I have more time for the shinystan r package?

```
rm(list=ls())  
load("modelo_Von_Bertanlanffy.RData")  
launch_shinystan(ajuste_Von_Bertanlanffy)
```

References

- ▶ Baptiste Auguie (2017). gridExtra: Miscellaneous Functions for “Grid” Graphics. R package version 2.3.
<https://CRAN.R-project.org/package=gridExtra>
- ▶ Jonah Gabry and Tristan Mahr (2018). bayesplot: Plotting for Bayesian Models. R package version 1.6.0.
<https://CRAN.R-project.org/package=bayesplot>
- ▶ Jonah Gabry (2018). shinystan: Interactive Visual and Numerical Diagnostics and Posterior Analysis for Bayesian Models. R package version 2.5.0.
<https://CRAN.R-project.org/package=shinystan>
- ▶ Stan Development Team (2018). RStan: the R interface to Stan. R package version 2.18.2. <http://mc-stan.org/>.

References

- ▶ https://mc-stan.org/docs/2_20/stan-users-guide/index.html
- ▶ https://mc-stan.org/docs/2_20/reference-manual/index.html
- ▶ https://mc-stan.org/docs/2_20/functions-reference/index.html